

Data ingestion in the modern enterprise

ETL is still used today for simple use cases, usually as part of a larger process. However, the complexity of today's data-intensive use cases precludes the use of ETL in most large enterprises.

For these reasons we refer collectively to the content of this part of the AI enterprise workflow as data ingestion rather than ETL.

The archaic process of ETL has given way to data ingestion approaches that emphasize the use of massive amounts of data *in place*, where analysis tools are used against data sources that have not been moved or transformed. This approach allows for much more flexibility since the transformation of data is isolated to each individual pipeline that is accessing a large data source. Different pipelines can be used to flexibly transform data from a large data store.